

Agent Model

Perceives its environment through sensors and acting upon that environment through actuators

- Percepts – Sensors
- Actions – Actuators
- Environment
- Performance Measure

Rational Agent

Select actions that maximises its (expected) utility
Percepts, Env, Action Space → Action Selected
I.e., Specified by an agent function $f: P \rightarrow A$

Rationality

- What is rational at a given time depends on PEAS:
- Performance measure
 - Prior Environment knowledge
 - Actions/Actuators
 - Percept sequence to date (Sensors)

Limitations:

- Percepts may not provide all the required info (Rationality ≠ omniscience)
- Actual outcome of actions may not be as expected (Rationality ≠ clairvoyant)

Environment Types

- **Fully observable (vs. partially observable):**
An agent's sensors give it access to the complete state of the environment at each point in time.
- **Deterministic (vs. stochastic):**
The next state of the environment is completely determined by the current state and the action executed by the agent.
 - If the environment is deterministic except for the actions of other agents, then the environment is **strategic**
- **Episodic (vs. sequential):**
The agent's experience is divided into atomic "episodes" (each episode consists of the agent perceiving and then performing a single action), and the choice of action in each episode depends only on the episode itself.
- **Static (vs. dynamic):** The environment is unchanged while an agent is deliberating.
 - (The environment is **semi-dynamic** if the environment itself does not change with the passage of time, but the agent's performance score does)
- **Discrete (vs. continuous):**
A limited number of distinct, clearly defined percepts and actions.
- **Single agent (vs. multiagent):**
An agent operating by itself in an environment

Search Problem Formulation

- **State space**, e.g. At(Arad), At(Bucharest)
- **Initial state**, e.g. At(Arad)
- **Actions**, set of actions given a specific state
 - **Transition model** e.g., Result(At(Arad),Go(Zerind)) → At(Zerind)
 - **Path cost** (additive), e.g., sum of distances, number of actions, etc
- **Goal test**, can be
 - Explicit, e.g. At(Bucharest)
 - Implicit, e.g. checkmate(x)

Search Problem Solution

- A **solution** is a sequence of actions from the initial state to a goal state (E.g., Arad → Sibiu → Fagarus → Bucharest)
- An **optimal solution** is a solution with the lowest path cost

General Search

- **Root** = Initial State, **Leaves** = Generated State
- **State** is a repr of a physical configuration
- **Node** is a data structure constituting **part of a search tree** (Comprises of state, parent, child, action path-cost, depth)
- **Expand** function creates new nodes
 - Uses Actions and Transition Model to create corresponding states


Search Strategies

Defined by picking the order of node expansion
Evaluated through:


- **Completeness** - find a solution if one exists?
- **Optimality** - least-cost solution?
- **Time complexity** - number of nodes generated/expanded
- **Space complexity** - maximum number of nodes in memory

Time and space complexity measured in terms of:

- **b** – max branching factor
- **d** – depth of least-cost solution
- **m** – max depth of state space

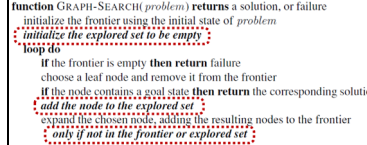


General Tree Search



Problem: Repeated states (Redundant paths can cause a tractable problem to become intractable)

General Graph Search



Types of Searches

- **Uninformed Search**
No extra info about states beyond that in the problem definition
- **Informed Search**
Uses problem-specific knowledge beyond the definition of the problem itself
- **Adversarial Search**
Used in multi-agent environment where the agent needs to consider the actions of other agents and how they affect its performance

Breadth-First Search

- General idea: Expand shallowest unexpanded node
Implementation: Use First-In First-Out (FIFO) queue
- **Completeness:** Yes (if b is finite)
 - **Optimality:** Yes (if cost=1 per step; Not optimal in general)
 - **Time complexity:** $1 + b + \dots + b^d = O(b^{d+1})$
 - **Space complexity:** $O(b^d)$, all node in memory

Problems: Memory Requirements & Execution Time

Uniform Cost Search

- General idea: Expand unexpanded node n with the lowest path cost g(n)
Implementation: Using a priority queue ordered by path cost g
- **Completeness:** Yes (if step cost > ε, some positive constant)
 - **Optimality:** Yes
 - **Time complexity:** $O(b^{C^*/\epsilon})$, where C^* = cost of optimal solution
 - **Space complexity:** $O(b^{C^*/\epsilon})$

Problems: Possible redundant searches

Depth-First Search

- General idea: Expand deepest unexpanded node
Implementation: Use Last-In First-Out (LIFO) queue
- **Completeness:** No (if m is infinite)
 - **Optimality:** No
 - **Time complexity:** $O(b^m)$, bad if m>d by a lot
 - **Space complexity:** $O(bm)$, linear space

Depth-Limited Search

- General idea: Expand deepest unexpanded node
Depth-First Search with predetermined depth limit l (Nodes at depth l have no child nodes & Solves infinite-path problem)
- **Completeness:** No (if $l < d$)
 - **Optimality:** No (if $l > d$)
 - **Time complexity:** $O(b^l)$
 - **Space complexity:** $O(bl)$

Iterative Deepening Search

- General idea: Use increasing Depth-Limited Search (DLS) to find the best depth limit l
- I.e., use DLS with depth limit 1. If no solution, then increase depth limit to 2. So on and so on, until solution is found
- Best of both Breadth-First Search and Depth-First Search
- **Completeness:** Yes
 - **Optimality:** Yes
 - **Time complexity:** $O(b^d)$
 - **Space complexity:** $O(bd)$

BFS vs DFS

- Use BFS when:
- Optimal solution is important
 - m is much greater than d
- Use DFS when:
- Space is important. DFS: $O(bm)$, BFS: $O(b^d)$

Heuristics

- The heuristic function h(n) is an estimate of how close a state n is to the goal state
- Informed search algorithms use heuristics to solve the search problem

Greedy Best-First Search

- General idea: Expand the node n with the lowest heuristic h(n)
Implementation: Use a priority queue ordered by heuristic h(n)
- **Completeness:** No (Can get stuck in loops, unless we keep track of repeated nodes)
 - **Optimality:** No
 - **Time complexity:** $O(b^m)$
 - **Space complexity:** $O(b^m)$, all nodes in mem

UCS vs G-BFS

- UCS is complete and optimal but may waste search in the wrong direction
- Greedy search generally in the correct direction but not complete or optimal
- Combine UCS & G-BFS → A* Search

A* Search

- General idea: Expand the node n that has incurred the least cost and is nearest to the goal state
Implementation: Using a priority queue ordered by eval. func. f(n)
- Evaluation function $f(n) = g(n) + h(n)$
 - Path cost g(n) = total path cost from start node to node n
 - Heuristic h(n) = estimated distance from node n to goal state
 - **Completeness:** Yes (if step cost > ε, some positive constant)
 - **Optimality:** Yes (If heuristics are admissible/consistent)
 - **Time complexity:** $O(b^{C^*/\epsilon})$, where C^* = cost of optimal solution
 - **Space complexity:** $O(b^{C^*/\epsilon})$

Applications:

- Path finding problems
- Video games
- Resource planning problems
- Robot motion planning

Heuristic Properties

- Admissibility:** A heuristic h(n) is admissible if $h(n) \leq h^*(n)$. For example:
- h(n) = estimated distance from node n to goal state
 - h*(n) = true cost from node n to goal state

- Consistent:** A heuristic h(n) is consistent if $h(n) \leq c(n, a, n') + h(n')$. For example:
- h(n) = estimated distance from node n to goal state G
 - h(n') = estimated distance from node n' to goal state G
 - c(n, a, n') = cost of getting from node n to n'

- Dominance:** A heuristic $h_2(n)$ dominates $h_1(n)$ if $h_2(n) \geq h_1(n)$, for all n.
- Only if both heuristics are admissible
 - A more dominant heuristic will be better for search (Potentially explore less branches)

Designing Heuristics

Admissible heuristics can be derived from the exact solution cost of a relaxed problem

Constraint Satisfaction Problems

- State**
- Defined by variables X_i that take on values from domain D_i
- Goal Test**
- A set of constraints C_i specifying allowable combinations of values for subsets of variables
- In contrast to standard search problems
- State is a "black box" - any old data structure that supports goal test, eval, successor

CSP comprises of:

- Finite set of variables $X = \{X_1, X_2, \dots, X_n\}$
- Non-empty domain D of k possible values for each variable D_i , where $D_i = \{v_1, \dots, v_k\}$
- Finite set of constraints $C = \{C_1, C_2, \dots, C_m\}$
- Each constraint C_i limits the values that variables can take, e.g., $V_1 \neq V_2$

Complete: Every variable is assigned
Consistent: Does not violate any constraint
CSP solution: Complete & Consistent assignment

Advantages of CSP

- **Formal representation language** that can be used to formalize many problem type
- Represent problem as a CSP and solve with general-purpose solver
- Can use **general-purpose solver**, which are more efficient than standard search
- Constraints allow us to focus the search to valid branches
- Branches that violate constraints are removed
- Non-trivial to do this for standard search (need manual selection of actions)

Constraint Graph

Nodes = Variables, Edges = Constraints

Variety of CSPs

- Discrete Variables
- Finite domains: $O(d^n)$ complete assignments for n variables, domain size d
 - Infinite domains: Integer, Strings, etc.
- Continuous Variables
- Time, float, etc.

Variety of Constraints

- **Unary:** Involve single variable
- **Binary:** Involve pairs of variables
- **Higher order:** Involve 3 or more variables
- **Preference** (Soft constraints)

CSPs as Standard Search

- Can be easily formulated (Initial State, Actions, Path Cost, Goal State)
- Sequence of actions do not matter, only the goal state (i.e., solution at depth n, use DFS)
- However, there are potentially $n!$ d^n leaves

Commutativity

- CSP variable assignments are commutative: i.e., Regardless of variable assignment order
- Only need to consider assignments to a single variable at each level/depth
- reduce from $n! d^n$ leaves to d^n leaves

Backtracking Search

- DFS for CSPs with single variable assignment
- Backtracking occurs when there are no legal values for a variable
- The basic uniformed algorithm for CSPs

General Purpose Methods

Can give huge gains in speed:

Minimum remaining values

- Choose the variable with the fewest legal values (i.e., the most constrained variable)

Degree Heuristic

- When multiple variables have the same MRV
- Choose the variable with the most constraints on remaining variables

Least constraining value

- Given a variable, choose the least constraining value
- The one that rules out the fewest values in the remaining variables

Forward Checking

- Keep track of remaining legal values for unassigned variables
- Terminate search when any variable has no legal values
- propagates information from assigned to unassigned variables
- Does not provide early detection for all failures
- Need to enforce constraints locally

Arc Consistency

- Simplest form of propagation make each arc consistent
- $X \rightarrow Y$ is consistent iff. for every value X there is some allowed y
- Arc consistency detects failure earlier than forward checking
- Can be run as a pre-processor or after each assignment
- Ordering of arcs do not matter
- Complexity of $O(n^2 d^3)$:
 - $O(n^2)$: Need to check for all edges, potentially n^2 edges
 - $O(d^2)$: For each edge, comparing their two domains
 - $O(d)$: Each variable change repropagate to neighbours, max d times

```
function AC-3(csp) returns the CSP, possibly with reduced domains
inputs: csp, a binary CSP with variables {X1, X2, ..., Xn}
local variables: queue, a queue of arcs, initially all the arcs in csp
while queue is not empty do
    (Xi, Xj) ← REMOVE-FIRST(queue)
    if REMOVE-INCONSISTENT-VALUES(Xi, Xj) then
        for each Xk in NEIGHBORS(Xj) do
            add (Xi, Xk) to queue
function REMOVE-INCONSISTENT-VALUES(Xi, Xj) returns true iff succeeds
removed ← false
for each x in DOMAIN[Xi] do
    if no value y in DOMAIN[Xj] allows (x,y) to satisfy the constraint Xi ↔ Xj
        then delete x from DOMAIN[Xi]; removed ← true
return removed
```

Problem Structure

- Suppose each subproblem has c variables out of n total
- Worst-case solution cost is $(\frac{n}{c} \times d^c)$

Tree-Structured CSPs

- if the constraint graph has no loops, the CSP can be solved in $O(nd^2)$ time
- 1. Choose a variable as root, order variables from root to leaves such that every node's parent precedes it in the ordering
- 2. For j from n down to 2, apply RemoveInconsistent(Parent(X_j), X_j)
- 3. For j from 1 to n, assign X_j consistently with Parent(X_j)

Nearly Tree-Structured CSPs

- Conditioning: instantiate a variable, prune its neighbours' domains
- Cutset conditioning: instantiate (in all ways) a set of variables such that the remaining constraint graph is a tree
- Cutset size $c \Rightarrow$ runtime $O(d^c \cdot (n - c)d^2)$, fast for small c

Representing Game as Search Problem

- Initial State
- Actions
- Terminal Test (Win/Lose/Draw)
- Utility Function (Numerical reward for the outcome):
E.g., Chess → +1, 0, -1
Poker → Cash win or lose
- Zero-sum: each player's utility for a state are equal and opposite

Minimax

- Perfect play for deterministic, perfect-information (fully observable) games
- Idea: choose moves with highest minimax value
 - best achievable payoff wrt best play
- **Completeness:** Yes, if tree is finite
- **Optimality:** Yes, against optimal opponent
- **Time complexity:** $O(b^m)$
- **Space complexity:** $O(bm)$

